

AI-Guided Memory Hierarchy Design For Edge LLM Inference

Hao Ding & Sichen Li

https://github.com/peylix/lm_memory_controller.git





Problem Statement & Motivation

Efficient memory hierarchy design is critical for enabling real-time, privacy-preserving LLM inference on edge devices under strict resource constraints.

- Edge LLM inference is memory-bound
- Limited on-chip SRAM, expensive off-chip DRAM
- Frequent movement of weights / activations / KV-cache
- Memory bandwidth and latency dominate performance & energy





Design Methodology

How We Used the Cognichip Platform: AI - Guided Design Workflow

- Started from natural-language descriptions of:
 - Edge LLM inference constraints
 - Memory-bound workload characteristics
- Used Cognichip AI to propose multiple design configurations, including:
 - Buffer sizes and memory partitioning
 - Tile granularity and access patterns
 - Prefetch and data-movement strategies
- Selected representative configurations for implementation



Design Methodology

How We Used the Cognichip Platform:

RTL & Simulation Development

- Used AI assistance to bootstrap the RTL framework
- Implemented a parameterized memory controller design
- Built a cycle-level simulation testbench to:
 - Drive tiled workloads
 - Model realistic memory behavior
 - Validate functional correctness

Analytical Performance Model

Purpose

Reduce the RTL design search space under limited project time.

Optimization Method

Evaluated design dimensions:

- Buffer size allocation
- Tile granularity
- Prefetch depth and overlap strategy

Multi-objective evaluation using Pareto frontier:

- Minimize DRAM accesses
- Minimize compute idle cycles
- Constrain SRAM footprint





Model Principle

The model estimates memory traffic, latency, and compute utilization across different buffer sizes, tile granularities, and prefetch strategies.

Outcome

We selected a small set of high-value configurations for RTL implementation and cycle-level simulation.

Analytical Performance Model



The optimal parameters are GEMM-specific tiling configurations selected per layer from the Pareto frontier, optimised primarily for the **prefill** stage.

Analytical Performance Model

Metric	Baseline	Optimized	Improvement
DRAM traffic	28.5 GB	14.4 GB	-49.3%
Compute utilization	61.7%	99.8%	+38.1 pp
Per-layer latency	3.65 ms	2.26 ms	1.62× faster
Full model latency	131.6 ms	81.4 ms	1.62× faster

By applying per-GEMM, prefill-focused tiling optimisation, we reduce DRAM traffic by **49.3%**, increase compute utilisation to near-peak (**99.8%**), and achieve a **1.62×** end-to-end latency reduction.




Architecture

We target a representative edge SoC configuration with **2 MB** on-chip SRAM (4×512 KB banks) and **50 GB/s LPDDR5** DRAM, running quantized (INT4/INT8) billion-parameter Transformer models.





RTL Control & Data flow Implementation



- Implements a centralized hardware control plane for tiled GEMM execution.
- Separates configuration, scheduling, arbitration, and prefetch into cycle-accurate RTL stages.
- Enforces execution correctness via hardware-gated start, busy tracking, and error handling.
- Overlaps computation and memory movement through ping-pong buffering and prefetch control.
- Exposes performance counters to support profiling and optimization of memory-bound execution.


Simulation results




8 Testbench files

50+ Tests pass


ensuring high coverage
and robust RTL validation.



Simulation results



Simulation results



Challenges & lessons learned

Challenges

- Large design space with expensive RTL simulation
- Memory behavior dominates performance and is hard to model accurately
- Coordinating parameters across RTL, testbench, and analytical model

Lessons Learned

- Early analytical modeling is critical for feasibility
- Clear module boundaries simplify iteration and debugging
- Performance counters are essential for understanding bottlenecks





Future work

- Explore adaptive or runtime-tunable memory management policies
- Extend the analytical model to larger LLM layers and multi-layer pipelines
- Run additional RTL simulations across more configurations



Thanks!

